Classification Algorithms

CSE347

Lauren Bright
Sophie Makaridi
Yanqi Liu

Choice of Datasets and Algorithms

Datasets:

Cho

1	1	-0.69	-0.96	-1.16	-0.66	-0.55	0.12	-1.07	-1.22	0.82	1.4	0.71	0.68	0.11	-0.04	0.19	0.82
2	1	-0.21	0.19	0.86	0.04	-0.35	-0.39	-0.51	-0.2	0.0	0.77	0.41	0.14	-0.45	-1.23	-0.325	0.0
3	1	-0.3	-0.56	-0.29	-0.5	-0.27	-0.29	-0.56	-1.04	0.32	0.9	0.45	0.17	0.164	-0.12	-0.16	0.67
4	1	0.07	0.26	-0.47	-0.68	-0.63	-0.39	0.07	0.79	0.58	0.31	-0.14	-0.29	-0.103	-0.2	-0.06	0.36
5	1	-1.04	0.13	0.51	-0.44	-0.88	-0.32	0.21	0.95	1.07	0.38	0.01	-0.13	-0.78	-0.13	0.092	0.0
6	1	-1.17	0.09	-0.52	-1.04	-1.16	-0.83	0.17	0.93	0.89	0.52	-0.24	-0.46	-0.215	0.2	0.91	0.68
7	1	-0.16	0.35	-0.13	-0.26	-0.4	-0.47	0.1	0.74	0.45	0.04	-0.3	-0.3	-0.118	-0.59	0.14	0.2

Cifar10

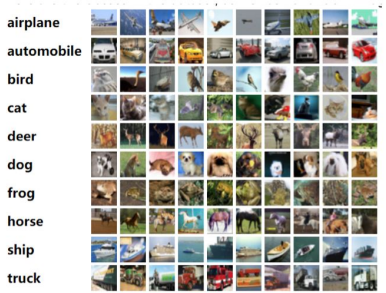

Algorithms:

Random Forest Classifier

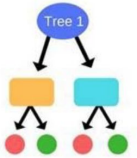

K-nearest neighbors

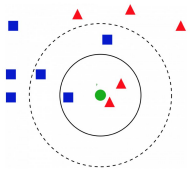

Support Vector Machine

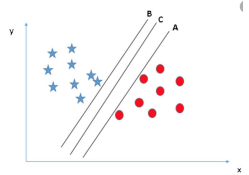

Data Processing and Classifying

Load Datasets from the drive

Cho: Removed ID and truth values. Removed outliers(-1)

Used MinMax method to map all values to [0,1] range.

Data and Labels were separated into two arrays

Configure Kfold. K=5 folds Set random_state seed

Used K-fold to get different combination of training/test sets

Initiate a model for the chosen algorithm and fit the model with training set

Predict label values using test sets

Evaluated performance using: Accuracy, AUC, Recall for each iteration and for all

Analyze Results

Algorithm Overview - Random Forest

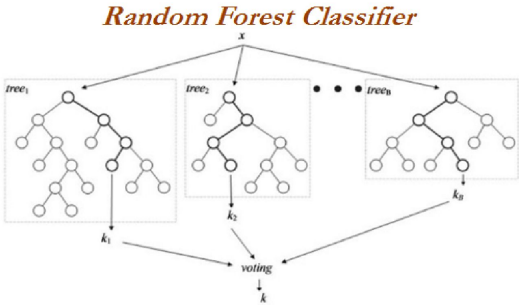

How it works:

- Random forest is an ensemble, tree-based algorithm
- The RF classifier generates a set of classification trees that each classify an object and “vote” for the class of the object
  - In this implementation, each tree outputted a probabilistic prediction, so not just a vote for one class, but different weights that could contribute to several classes
- The overall forest then considers all votes and chooses a final classification for the object
- Having several trees cuts down on variance and overfitting that can occur with just one tree
- The main parameter is **n_estimators**, which determines the number of trees in the forest
  - n_estimators = 100 was chosen with the same seed for each run
  - Choosing a higher value for n_estimators could increase accuracy, but greatly increases run time

Image from: https://www.mygreatlearning.com/blog/random-forest-algorithm/

Algorithm Overview - K-Nearest Neighbor

How it works:

- The k-nearest neighbors (KNN) algorithm is a simple, easy-to-implement supervised machine learning algorithm.
- Assumes that similar things exist in close proximity.
- Stores all available cases and classifies new cases based on a similarity measure.
  - Distance function - Euclidean distance
- A case is classified by a vote of its neighbors.
  - Majority voting vs. weighted voting

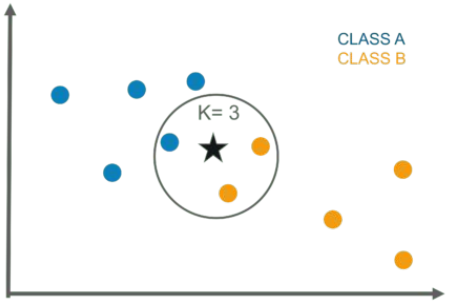

Algorithm Overview - Support Vector Machine

How it Works:

- find a hyperplane in an N-dimensional space to separate classes
- Out of all possible hyperplanes maximize the margin.
- Uses Lagrangian to solve the optimization problem
- When not linearly separable, use Kernel functions : linear, poly, rbg, sigmoid, precomputed, callable

- I used the Radial Basis Function (RBF).
  - Creates additional features to increase dimensions
  - Points that were hard to classify linearly, become easily separable in higher dimensions.

Setbacks - difficulties

- K-nearest neighbors and Support vector machine are time inefficient on large datasets
  - Cifar10 took hours to finish
  - Made it hard to test with different parameters
  - A way to go around was to test on a smaller part of Cifar
- Working with datasets
  - Datasets were fairly different (Image vs. numerical)
  - Understanding datasets, features, type, size
- Detecting Outliers/noise
  - Hard to see for cho
  - Easier for Cifar since images are easier to visualize

Comparing Results

SVM (Cho):

Average	0.7695	0.8417	0.7658
Standard Deviation	0.0510	0.0461	0.0348

SVM (CIFAR):

Average	0.5455	0.7082	0.5470
Standard Deviation	0.0510	0.04611	0.0349

KNN (Cho):

Average	0.7383	0.8282	0.7248
Standard deviation	0.0213	0.0122	0.0198

KNN (CIFAR):

Average	0.3388	0.6327	0.3389
Standard deviation	0.0043	0.0016	0.0027

Random Forest (Cho):

Average	0.7305	0.8246	0.7193
Standard Deviation	0.0295	0.0222	0.0378

Random Forest (CIFAR):

Average	0.4641	0.7023	0.46412
Standard Deviation	0.0046	0.0025	0.0045

Observations and takeaways

Observations:

- Large vs. small datasets with different algorithms
- Different types of data perform differently
  - Image vs. numerical values
  - Question of quality of data

Takeaways:

- Parameter tuning is difficult and time consuming
  - Many parameters to consider
- Different algorithms work better for certain tasks